

WHICH CONDITIONING RECORDING DOES A VOICE CLONE FOLLOW? A BALANCED CROSS-MICROPHONE CROSSOVER

Rafaello Virgilli Lucas Alcântara Souza Alexandre Costa Ferro Filho Daniel Tunnermann
Gustavo dos Reis Oliveira Arlindo Rodrigues Galvão Filho Anderson da Silva Soares

Advanced Knowledge Center in Immersive Technologies (AKCIT)
Federal University of Goiás (UFG), Goiânia, GO, Brazil

ABSTRACT

Which of two same-speaker recordings supplied a voice clone’s conditioning input? In a pre-specified complete crossover over 54 VCTK speakers selected after earlier ECAPA-based screening, we clone from each of two utterances and compare each clone with both, recorded through the other microphone; four systems and four texts give 3,456 clones. Generation and attribution use opposite microphones, so speaker identity, generated text, exact waveform and a shared microphone fingerprint cannot decide the label. With fixed speaker-verification encoders, speaker-averaged two-alternative attribution accuracy for mic1→mic2 is .896 [.869,.921] for ECAPA and .622 for WavLM; reversing the microphones gives .907 [.881,.931] and .620. When both candidates are themselves clones from a different system and text, ECAPA attribution remains .805 [.780,.830] and .805 [.780,.831]; all 108 VCTK speakers with paired captures give pooled ECAPA .909/.909. Brackets are 95% whole-speaker bootstrap intervals. The main evaluation compares two same-speaker candidates known to include the source event; prompt transformations do not identify or exclude the carrier.

Index Terms— voice cloning, source attribution, audio provenance, speaker verification

1. INTRODUCTION

Zero-shot speech generation conditions on a short reference recording. When that recording was used without consent, identifying the generator or cloned speaker does not answer the narrower provenance question: *which recording was used?* Existing source tracing attributes a synthesis pipeline [1, 2], and clone attribution identifies an enrolled speaker [3, 4]. Recording-level evidence would instead connect an output to one conditioning event among other recordings of the same voice.

That target is easy to confound: a clone and candidate share speaker identity, one candidate may be intrinsically closer to every clone, recording session, duration or microphone can reveal the answer, and generated content may overlap. Prior prompt-transfer work shows preservation of prosody and acoustic environment [5, 6], while speaker embeddings encode session and channel [7], so high retrieval from a same-speaker pool does not by itself show that the target follows the recording used for conditioning.

We turn the question into a balanced intervention. For each speaker we clone from recording event A and event B and test both signed comparisons, with prompt and candidate paired captures from different microphones and systems and texts fully crossed. ECAPA-TDNN [8] (SpeechBrain `spkrec-ecapa-voxceleb`) is the fixed primary encoder, with a speaker-verification-fine-tuned

WavLM [9] (`microsoft/wavlm-base-plus-sv`) for corroboration. The experiment uses a 54-speaker complement inherited from ECAPA-informed development rosters; all speakers in it are retained, and metadata alone determines pair selection without clone outcomes.

The complete A/B crossover, whose analysis plan and decision rule were fixed before outcomes were inspected, tests whether a clone follows its conditioning recording after a pickup-path change over 3,456 clones, 54 speakers, four systems and two fixed representations. Attribution then persists when both candidates are themselves clones made with a different generator and a different output text from the query: ECAPA reaches .805 in both microphone directions, with prompt microphones crossed and speaker identity held fixed. The boundary is explicit: recording-event attribution is consistent in a two-candidate closed set, but the experiment identifies neither the physical carrier nor deployable recording-presence detection.

2. RELATED WORK

Attribution and provenance. Generator tracing, shared-generator verification and speaker attribution operate above the individual recording [1, 2, 3, 4]. Recent source-verification and passive-fingerprint studies distinguish generating systems and examine attribution cues [10, 11, 12]; our label instead identifies the conditioning event within one speaker. Watermarking requires instrumentation before misuse [13]; VoiceMark attributes cloned speech through a watermark embedded in its inference-time reference [14], whereas our candidates are unmodified recordings. Training-set membership inference asks whether audio entered training [15], and VoxGuard tests enrolled-speaker and attribute privacy [16]; neither asks which recording was the inference-time prompt. We found no prior controlled study that intervenes on the prompt recording within speaker while changing the candidate pickup path.

Prompt transfer and nuisance information. Training-free generator attribution uses frozen representations [17, 18], and recording-related nuisance in speaker embeddings motivates cross-recording sampling [19]. Prompt-conditioned systems transfer prosody and environment [5, 6], although cloning also alters style systematically [20]; prompt cosine is a standard synthesis metric [21]. Exact prosody cloning transfers reference prosody [22]. Source-voiceprint tracing recovers identity from converted audio [23], while rank/local-disclosure work treats linkage without global thresholds [24, 25]. None intervenes on which of two same-speaker recording events conditions a clone. Concurrent work separates prompt environment from identity [26]. These works motivate recording-level leakage but do not identify the chosen event. We target *A versus B*,

Primary: opposite-pickup attribution

$$A_{m1} \rightarrow c_{A,m1} \rightarrow \{A_{m2}, B_{m2}\} \quad \text{and} \\ B_{m1} \rightarrow c_{B,m1} \rightarrow \{A_{m2}, B_{m2}\}$$

Correct labels are reversed across arms; speaker, system and generated text are held fixed. The direction-reversal analysis swaps $m1$ and $m2$.

Fig. 1. Complete two-arm crossover. Generation never consumes the microphone used for the attribution candidates.

not generic clone–prompt resemblance.

3. CROSS-MICROPHONE CROSSOVER

3.1. Task and design

For speaker s , let recording events A and B each have paired VCTK mic1/mic2 captures [27]: a DPA 4035 omni and a Sennheiser MKH 800. That the two are genuinely distinct signals is the design’s load-bearing assumption. We tested it rather than assuming it. Raw correlation cannot settle the question. A delayed copy can score near zero, while genuine pairs reach .92 median once delay-aligned. We therefore use the normalized least-squares residual $\min_{g, |\tau| \leq 50 \text{ ms}} \|y - g T_\tau x\|_2 / \|y\|_2$ over a scalar gain g and a zero-padded delay T_τ , where x and y are the mic1 and mic2 captures; its .00001 duplicate bar is a numerical tolerance for these transforms, not a universal perceptual threshold. Across all 108 event captures, the residual is .38 median (range .21–.70). Exact, gain-scaled and delayed injections on every source, including both delay-window boundaries, are at most .000000026; no pair is byte-identical. A clone $c_{A,m1}$ is conditioned on event A ’s mic1 capture. The primary comparison asks whether a fixed encoder f gives

$$I[\cos(f(c_{A,m1}), f(A_{m2})) > \cos(f(c_{A,m1}), f(B_{m2}))], \quad (1)$$

with ties worth 1/2. We also clone from B and reverse the correct label. Thus an encoder’s unconditional preference for A cannot create a positive effect. The predeclared direction reversal swaps microphones (mic2 prompt, mic1 candidates) and cannot rescue a failed primary direction.

With $q(u, v) = \cos(f(u), f(v))$ and $d_X = q(c_X, A) - q(c_X, B)$ on opposite-microphone candidates, a cell contributes $\{I[d_A > 0] + I[d_B < 0]\}/2$. Clone-independent candidate bias therefore yields chance; ties receive 1/2.

The threat model is deliberately narrow. The analyst has a clone and two same-speaker candidate recordings; exactly one is the other-microphone capture of the event that supplied the conditioning recording. The objective is source attribution, not deciding whether an arbitrary recording was used. Encoder parameters, score direction and the tie rule are fixed before evaluation. The analyst does not train on these speakers or systems.

The design addresses five shortcuts and names the residual of each: A/B within speaker removes speaker identity, both signed arms remove a fixed candidate bias and leave arm interactions, four texts shared across A and B remove generated content and stay fixed rather than sampled, the opposite microphone removes the exact waveform and device and leaves event acoustics, and metadata-matched pairs limit duration and length cues but not other local ones.

3.2. Frozen roster and generation

VCTK v0.92 documents DPA 4035 and Sennheiser MKH 800 microphones in a fixed hemi-anechoic setup [27]. Its public description does not establish acquisition timing, so our claim requires paired

captures, not simultaneity. The frozen manifest independently supports the pairing: all 108 selected event/microphone pairs have identical frame counts and sampling rates.

The 54-speaker roster is the exact complement of two earlier development subsets. The first applied an ECAPA pair-separation screen to 105 speakers (pair cosine below that speaker’s 90th percentile over pairs of duration-eligible utterances) and retained the lexicographically first 30 passing speakers. The second retained all 24 remaining speakers with one eligible pair at ECAPA cosine distance $\leq .20$ and another at $\geq .38$, pair-mean durations within .50 s. The roster is therefore disjoint from those two subsets, although its speakers appeared in an earlier exploratory screen. It is blind to this experiment’s clone outcomes, but not embedding-naïve or population-sampled. Within that fixed complement, all 54 speakers remain. The A/B-pair rule below never inspects an embedding, score or clone outcome.

For each speaker, selection enumerates 4–10 s paired utterances and excludes the conditioning utterance used for that speaker in the earlier exploratory screen (listed in the public selection manifest). It requires an A/B duration gap of at most .25 s and a relative seconds-per-UTF-8-byte gap $|r_A - r_B| / [(r_A + r_B)/2]$ of at most 5% on *both* microphones. A deterministic metadata-only rule then takes the lexicographic minimum of the largest tolerance-normalized gap, their sum, the largest rate gap and the largest duration gap; no embedding, score or clone outcome enters within-roster pair selection. Transcripts A and B differ. Generation instead uses four fixed texts that are disjoint from both candidates and fully crossed across arms. These four texts were also used in our earlier experiments; they are fixed experimental factors. A post-hoc ASR screen (Whisper-large-v3-turbo) gives median WER .04 against the requested text; no clone (0 of 3,456) was flagged for four or more consecutive recognized words shared with its conditioning transcript after excluding sequences present in the requested text. This diagnostic does not establish the absence of reference-speech repetition. The selected roster is substantially tighter than the eligibility limits. Its worst duration gap is .112 s, and its worst relative duration-per-byte gap is 2.57% across either microphone.

We evaluate F5-TTS (f5-tts 1.1.22) [28], XTTS-v2 (coqui-tts 0.25.3) [29], CosyVoice 2 [30] and Seed-VC [31]; exact checkpoint revisions are listed in the public artifact. Seed-VC converts one fixed LibriTTS-R [32] source utterance per generation text, shared across targets. The full design is 54 speakers $\times$ 2 events $\times$ 2 prompt microphones $\times$ 4 texts $\times$ 4 systems = 3,456 clones. Exact checkpoints, repositories, environments, files and per-clone ancestry ledgers are hash-pinned. Generation seeds are fixed by output-text index and shared across arms and speakers, so the intervals also condition on these realizations.

3.3. Estimand and decision criteria

Each speaker contributes 32 dependent comparisons per direction (two arms, four systems, four texts). We average within speaker, then weight all 54 speakers equally. Percentile intervals resample whole speaker vectors 100,000 times with a frozen seed; they quantify stability to speaker composition in this fixed eligible roster, not population-coverage confidence. Systems and texts are fixed crossed factors, and system-specific points are descriptive.

The primary success criteria, fixed before outcomes were inspected, were ECAPA accuracy $\geq .80$ with lower stability-interval endpoint $> .70$, and WavLM lower endpoint $> .50$, an author-declared large-effect criterion, not a deployment-utility claim. Both reversed-microphone estimates are reported; corroboration after re-

versal requires both lower endpoints above .50. Diagnostics cannot rescue any criterion.

We fixed the roster, analysis settings and success criteria before inspecting outcomes. A second, independent recomputation rehashed 3,456 clones and ledgers plus 216 candidates and reproduced every result field to 2×10^{-15} . The released artifact contains per-clone cosine scores with candidate identities and an independent verifier that reconstructs the 3,456-row identity grid, bootstrap and verdict without audio or model weights.

4. RESULTS

4.1. The conditioning event survives pickup change

Table 1 reports the predeclared directions. ECAPA attributes the conditioning event in 89.6% of primary comparisons and 90.7% in reverse. It therefore meets the .80/.70 primary bar. WavLM is much weaker but directionally consistent at 62.2%/62.0%. All five predeclared criteria pass. The cross-microphone estimates are .896 [.869,.921] and .907 [.881,.931] for ECAPA and .622 [.593,.652] and .620 [.589,.653] for WavLM: the four lower endpoints exceed .50, the primary ECAPA point reaches .80, and its lower endpoint .869 exceeds .70. The ECAPA effect meets the frozen large-effect rule, and WavLM corroborates its direction.

Clone candidates. A stricter comparison, whose grid and decision rule were fixed before analysis, replaces both real candidates by clones: the query clone is compared with two clones of the same speaker made by a different system from a different text, conditioned on events A and B through the opposite microphone (288 comparisons per speaker and direction). The correct candidate shares only the conditioning event. Attribution remains .805 [.780,.830] and .805 [.780,.831] for ECAPA and .610 [.595,.626] and .612 [.593,.632] for WavLM (Table 1, Clone-to-clone column), so the event information survives a change of generator, output text, microphone and candidate medium.

Other rosters. A replication with the decision rule fixed before analysis re-runs the crossover on a second utterance pair per speaker, chosen by the same metadata-only rule after excluding the original pair (53 speakers; one admits no second pair; 3,392 clones). The rule passes again: ECAPA .906 [.878,.931] and .909 [.883,.933], WavLM .666 [.636,.696] and .644 [.616,.673] (Table 1). Observed per-speaker ECAPA accuracies have Pearson correlations .061/.291 between the two utterance pairs; this descriptive comparison does not establish the stability of an underlying speaker effect.

The same protocol applied to the 54 earlier-development speakers, whose historical pairs passed ECAPA-based selection criteria, gives ECAPA .922 [.895,.946] and .911 [.886,.935] and WavLM .656 and .630; no selected pair repeats an earlier development pair, and two speakers share one utterance with theirs. These ECAPA points exceed the complement’s .896/.907, a descriptive comparison that does not isolate an effect of earlier screening. Pooling their speaker means with the original 54 gives ECAPA .909 [.890,.927] and .909 [.891,.926] and WavLM .639 [.617,.661] and .625 [.602,.648] over the complete 108-speaker paired-capture roster, with no speakers excluded by the earlier screen.

Same-microphone diagnostics, never verdict inputs, are only modestly higher (ECAPA .938/.930, WavLM .631/.659, primary/reverse): exact pickup identity is unnecessary, though paired captures of one event do not establish microphone invariance.

Figure 2 exposes the heterogeneity hidden by pooled points. All 54 primary ECAPA speaker means exceed .50; in reverse, 53 exceed it and one equals it. WavLM is not uniformly positive: 47 primary

Table 1. Speaker-averaged two-alternative attribution accuracy; P/R: primary/reverse; stability intervals in §4.1. Comparisons per direction: 1,728 cross-mic, 15,552 clone-to-clone, 1,696 fresh pairs (53 speakers), 3,456 full roster (108 speakers). Clone-to-clone candidates come from a different system and text; the full roster pools the 54 earlier-development speakers with these.

Encoder, direction	Cross-mic	Clone-to-clone	Fresh pairs	Full roster
ECAPA, P	.896	.805	.906	.909
ECAPA, R	.907	.805	.909	.909
WavLM, P	.622	.610	.666	.639
WavLM, R	.620	.612	.644	.625




Fig. 2. Fixed-roster per-speaker attribution accuracy (54 dots/group; 32 comparisons per dot). Diamonds and bars are pooled points and 95% speaker-resampling stability intervals; the dashed line is chance. P/R: primary/reverse.

and 43 reverse speaker means exceed .50, so its role is aggregate corroboration, not a per-speaker guarantee. Whole-speaker resampling preserves this heterogeneity rather than treating 1,728 comparisons per direction as independent.

4.2. Breadth across fixed systems

Descriptive system points, primary/reverse for F5-TTS, XTTS-v2, CosyVoice 2 and Seed-VC, are .942/.942, .833/.822, .910/.942 and .898/.924 for ECAPA and .618/.657, .572/.574, .671/.657 and .627/.593 for WavLM. ECAPA exceeds .82 throughout, WavLM spans .572–.671 with weaker XTTS-v2 cells, and there are no simultaneous per-system intervals, so the breadth claim is limited to this tested design.

The ECAPA–WavLM gap rules out near-ceiling magnitude for every representation, while WavLM’s positive lower endpoints oppose ECAPA-weight specificity; both are speaker-verification embeddings, so a shared cue remains possible. Mean correct-minus-incorrect cosine margins are positive but small and descriptive only: ECAPA .079/.081, WavLM .011/.013 for primary/reverse.

4.3. Descriptive sensitivity analyses

Stratifying the inherited ECAPA-informed complement by the earlier pair-separation screen, ECAPA is .908/.907 among the 36 pass-but-not-retained speakers, .879/.915 among the 15 fail speakers and .833/.875 among the three absent from that record, with WavLM at .609/.596, .654/.685 and .625/.583; no stratum carries a confirmatory claim. This post-hoc check cannot create an embedding-naïve sample or alter the pooled verdict.

The crossover, the clone-candidate comparison and the fresh-pair replication have decision rules fixed before analysis. The additional-cohort evaluation and its pooling were specified before generating those clones, after the original and earlier development results were

known. The remaining analyses are post-hoc and descriptive. Two such extensions bound the claim. Enlarging the candidate set with metadata-matched captures of the same speaker through the opposite microphone, among the 22 speakers with enough eligible captures, ECAPA rank-1 accuracy decreases from .935/.913 with two candidates to .635 [.545,.720]/.587 [.510,.659] with sixteen (704 comparisons per direction; sixteen-candidate chance .0625). Resynthesizing each primary-direction clone with a different system and a different text, the second-generation clone is still attributed to the original event at .796 [.759,.832] (ECAPA) and .596 [.568,.626] (WavLM).

Arm-separated points also oppose a one-candidate shortcut. For A/B, ECAPA is .889/.903 in primary and .892/.922 in reverse; WavLM is .588/.656 and .660/.581, respectively. Thus WavLM’s arm asymmetry reverses with microphone direction. Descriptively, all 32 system×direction×encoder×arm cells are above .50 (minimum .523). No cell has a separate interval or multiplicity-adjusted test.

4.4. What changed relative to earlier evidence

An earlier exploratory two-reference crossover (30 speakers, 960 trials) obtained .806 [.757,.854] on one pickup path with favorable reference pairs and missed its frozen .90 point bar. The present experiment is a separate test with its own frozen .80/.70 ECAPA bar; passing it does not retroactively change the earlier .90 miss.

5. BOUNDARY AND LIMITATIONS

Attribution is not presence detection. The crossover is a closed-set two-alternative question known to contain the source. Table 2 sweeps one common threshold per encoder over the same scores, pooled over systems and arms, as one-to-one verification within speaker. At the .01 prior, minimum DCF sits at or near the reject-all cost of 1.000 (whole-speaker resampling gives ECAPA upper endpoints .989/.989); at the .5 prior of the two-candidate threat model it is .67/.68 and .88/.89. This post-hoc check is descriptive: a score can order A above B while no global threshold separates the two.

We therefore restrict investigative triage to a known-positive two-recording set for human provenance review: the score prioritizes candidates, cannot decide whether an arbitrary queried recording was present, and is not an automated accusation.

The carrier is unidentified. Opposite microphones exclude an exact waveform and shared device fingerprint, while A/B balancing excludes a fixed candidate preference, and metadata matching sharply limits duration shortcuts. Three separate post-hoc prompt transformations, synthesized with the same systems and scored against unmodified candidates (primary direction), test robustness rather than mechanism: ECAPA accuracy was .896 after total-duration equalization, .893 after smoothed long-term spectral equalization and .840 [.804,.874] after pitch/energy modification (.866

when the candidates receive the same transform), against .896 for the original prompts. These transformations do not isolate causal contributions: relative timing and residual spectral information may remain, and the pitch/energy comparison includes resynthesis effects. Post-hoc alternative readouts on the same grid remain below ECAPA: a long-term spectral and cepstral average reaches .548/.549, and mean-pooled WavLM layer 0, centred on the speaker’s real-capture mean, reaches .633/.631 against .622/.620 for WavLM-SV. The carrier therefore stays bundled in event-level prosody, articulation, background acoustics, paired-capture room state and their interactions with model conditioning.

Scope. The primary result concerns four fixed checkpoints, four fixed texts, read English speech and one inherited 54-speaker VCTK complement, disjoint from the earlier development rosters but ECAPA-informed through their composition. The added evaluation extends the crossover to all 108 speakers with paired captures, removing screen-based speaker exclusion from the pooled estimate; prior ECAPA-informed development remains, and these data are not globally unseen or population-sampled. Checkpoint training corpora are not documented at recording level, so VCTK exposure cannot be excluded. The intervals do not sample systems, texts or a speaker population. Spontaneous conversational speech, noisy field recordings, other pickup paths, candidate sets beyond sixteen and open-set search remain untested; larger candidate sets were examined only post hoc on eligible subsets.

Release and misuse. Code, score-level data, provenance metadata and aggregate/per-speaker results are at <https://github.com/rvirgilli/voice-clone-cross-microphone-crossover/tree/f2-icassp2027-rc3>; model weights, source audio and playable impersonations are not redistributed. Recording attribution can assist provenance audits, but it can also deanonymize speakers, and licensing alone does not settle consent.

6. CONCLUSION

In this fixed roster, clones preferentially match the other-microphone capture of their conditioning event across the two tested microphone paths. ECAPA meets the frozen large-effect criterion and WavLM shows a smaller effect in the same direction, and that direction persists when the microphone roles are reversed, when both candidates are clones from another system and text (.805), on a second utterance pair per speaker (.906/.909), over the complete 108-speaker paired-capture roster (.909/.909), and after a second synthesis (.796), all ECAPA points. The primary and reverse pooled estimates weight the four fixed systems equally. This is a large closed-set signal between two tested microphone paths; its physical carrier remains unidentified, and the experiment does not solve open-set recording-presence detection.

7. ACKNOWLEDGMENT

Funded by the project Research and Development of Algorithms for Construction of Digital Human Technological Components, supported by the Advanced Knowledge Center in Immersive Technologies (AKCIT), with financial resources from the PPI IoT of the MCTI grant 057/2023, signed with EMBRAPPI. The authors declare no competing interests.

During preparation, the authors used a large language model (Anthropic Claude) to help write and edit the text throughout the manuscript, the LaTeX layout and the supporting code. The authors reviewed all assisted content and take full responsibility.

Table 2. Within-speaker one-to-one verification, one global threshold. Target: clone versus its own event; non-target: clone versus the speaker’s other event; 1,728 each per direction. $C_{miss} = C_{fa} = 1$.

Encoder	Dir.	EER (fraction)	minDCF prior .01	minDCF prior .5
ECAPA	P	.337	.980	.667
ECAPA	R	.343	.981	.678
WavLM	P	.447	1.000	.884
WavLM	R	.451	.998	.889

8. COMPLIANCE WITH ETHICAL STANDARDS

This study used only previously collected, publicly available recordings, under the licence and terms of use of the VCTK corpus, for non-commercial academic research. It involved no new recording and no interaction with human participants, and required no ethical approval. Synthetic speech is not redistributed in playable form.

9. REFERENCES

- [1] Anton Firc et al., “STOPA: A database of systematic variation of deepfake audio for open-set source tracing and attribution,” in *Proc. Interspeech*, 2025, arXiv:2505.19644.
- [2] Viola Negroni et al., “Forensic similarity for speech deepfakes,” in *Proc. ACM IH&MMSec*, 2026, arXiv:2510.02864.
- [3] Shuhei Kato, “A geometry-limited identification floor and its consequences for voice-clone attribution in professional voice actors,” arXiv:2607.15694, 2026.
- [4] Candice R. Gerstner, “Audio avatar fingerprinting: An approach for authorized use of voice cloning in the era of synthetic audio,” arXiv:2603.20165, 2026.
- [5] Haitao Li, Chunxiang Jin, Chenglin Li, Wenhao Guan, Zhengxing Huang, and Xie Chen, “ReStyle-TTS: Relative and continuous style control for zero-shot speech synthesis,” in *Findings of the ACL*, 2026, pp. 9257–9269.
- [6] Xiaofei Wang, Sefik Emre Eskimez, Manthan Thakker, Hemin Yang, Zirun Zhu, Min Tang, Yufei Xia, Jinzhu Li, Sheng Zhao, Jinyu Li, and Naoyuki Kanda, “An investigation of noise robustness for flow-matching-based zero-shot TTS,” in *Proc. Interspeech*, 2024, arXiv:2406.05699.
- [7] Desh Raj, David Snyder, Daniel Povey, and Sanjeev Khudanpur, “Probing the information encoded in x-vectors,” in *Proc. IEEE ASRU*, 2019, arXiv:1909.06351.
- [8] Brecht Desplanques, Jenthe Thienpondt, and Kris Demuynck, “ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification,” in *Proc. Interspeech*, 2020.
- [9] Sanyuan Chen et al., “WavLM: Large-scale self-supervised pre-training for full stack speech processing,” *IEEE J. Selected Topics in Signal Processing*, vol. 16, no. 6, pp. 1505–1518, 2022.
- [10] Pierre Falez, Tony Marteau, Damien Lolive, and Arnaud Delhay, “Audio deepfake source tracing using multi-attribute open-set identification and verification,” in *Proc. Interspeech*, 2025, pp. 1528–1532.
- [11] Viola Negroni, Davide Salvi, Paolo Bestagini, and Stefano Tubaro, “Source verification for speech deepfakes,” in *Proc. Interspeech*, 2025, pp. 1548–1552.
- [12] Yupei Li, Qiyang Sun, Emmanouil Benetos, Berrak Sisman, and Björn Schuller, “Perceptible or not? Diagnosing passive fingerprints for speech deepfake attribution,” arXiv:2609.00765, 2026.
- [13] Yunming Zhang et al., “VocaLock: Watermark-based detection of zero-shot voice conversion manipulation and timbre attribution,” *IEEE Trans. Information Forensics and Security*, 2026.
- [14] Haiyun Li, Zhiyong Wu, Xiaofeng Xie, Jingran Xie, Yaoxun Xu, and Hanyang Peng, “VoiceMark: Zero-shot voice cloning-resistant watermarking approach leveraging speaker-specific latents,” in *Proc. Interspeech*, 2025, pp. 5108–5112.
- [15] Thomas Sesmat, Gabriel Meseguer-Brocal, and Geoffroy Peeters, “Where flow matching leaks: Characterising membership signals along the interpolation path,” in *Proc. ICML*, 2026, arXiv:2606.07271.
- [16] Efthymios Tsaprazlis et al., “VoxGuard: Evaluating user and attribute privacy in speech via membership inference attacks,” in *Proc. ICASSP*, 2026.
- [17] Adriana Stan, David Combei, Dan Oneata, and Horia Cucu, “TADA: Training-free attribution and out-of-domain detection of audio deepfakes,” in *Proc. Interspeech*, 2025, pp. 1543–1547.
- [18] Gabriel Pirlogeanu, Adriana Stan, and Horia Cucu, “Understanding the strengths and weaknesses of SSL models for audio deepfake model attribution,” in *Proc. IEEE ICASSP*, 2026, arXiv:2603.13488.
- [19] Theo Lepage and Reda Dehak, “SSPS: Self-supervised positive sampling for robust self-supervised speaker verification,” in *Proc. Interspeech*, 2025, pp. 1098–1102.
- [20] Kaitlyn Zhou, Federico Bianchi, Martijn Bartelds, Anna Pot, Yongchan Kwon, and James Zou, “Voice “cloning” is style transfer,” arXiv:2605.16578, 2026.
- [21] Chengyi Wang et al., “Neural codec language models are zero-shot text to speech synthesizers,” arXiv:2301.02111, 2023, VALL-E; SIM-o as a standard zero-shot TTS evaluation metric.
- [22] Florian Lux, Julia Koch, and Ngoc Thang Vu, “Exact prosody cloning in zero-shot multispeaker text-to-speech,” in *Proc. IEEE SLT*, 2022, arXiv:2206.12229.
- [23] Jiangyi Deng, Yanjiao Chen, Yinan Zhong, Qianhao Miao, Xueluan Gong, and Wenyuan Xu, “Catch you and I can: revealing source voiceprint against voice conversion,” in *Proc. USENIX Security*, 2023.
- [24] Tom Bäckström, Mohammad Hassan Vali, My Nguyen, and Silas Rech, “Privacy disclosure of similarity rank in speech and language processing,” *IEEE Trans. Audio, Speech, Lang. Process.*, 2026, arXiv:2508.05250.
- [25] Dāvis Šterns et al., “Goodbye equal error rate, hello local information disclosure: Evaluating voice anonymisation against 1-to-N linkage threats,” arXiv:2607.06259, 2026.
- [26] Ye-Xin Lu et al., “Towards real-world environment-aware zero-shot text-to-speech synthesis via disentangled audio infilling,” arXiv:2608.03011, 2026.
- [27] Junichi Yamagishi, Christophe Veaux, and Kirsten MacDonalld, “CSTR VCTK Corpus: English multi-speaker corpus for CSTR voice cloning toolkit (version 0.92),” 2019, University of Edinburgh, Centre for Speech Technology Research, sound dataset.
- [28] Yushen Chen et al., “F5-TTS: A fairytaler that fakes fluent and faithful speech with flow matching,” arXiv:2410.06885, 2024.
- [29] Edresson Casanova et al., “XTTS: A massively multilingual zero-shot text-to-speech model,” in *Proc. Interspeech*, 2024, arXiv:2406.04904.
- [30] Zhihao Du et al., “CosyVoice 2: Scalable streaming speech synthesis with large language models,” arXiv:2412.10117, 2024.
- [31] Songting Liu, “Zero-shot voice conversion with diffusion transformers,” arXiv:2411.09943, 2024, Seed-VC.
- [32] Yuma Koizumi et al., “LibriTTS-R: A restored multi-speaker text-to-speech corpus,” in *Proc. Interspeech*, 2023, arXiv:2305.18802.